

?Get ready for a party?: Exploring smarter smart spaces with help from large language

models Evan King, Haoxiang Yu, Sangsu Lee, Christine Julien The University of Texas at

Austin {e.king, hxyu, sethlee, c.julien}@utexas.edu 3202 raM 42 ]CH.sc[ GPT models can

infer meaning behind ambiguous user commands and control smart home devices in

response. When told

to ?setupforaparty?,GPT-3producesvalidJSONthatsetsthelightstoacolorloopandturnsonthestereoformusicplay

right response to someone who says ?get ready I. INTRODUCTION for a party? is deeply

influenced by meaning and context. For a smart home assistant (e.g., Google Home), the

ideal response An exciting prospect of smart homes at their advent

was mightbetosurveytheavailabledevicesinthehomeandchange the potential to reduce user

burden by providing seamless, their state to create a festive atmosphere. Current

practical unobtrusive,and?smart?interfacestoeverydaydevices.While systems cannot service

such requests since they require the ability to (1) infer meaning behind an abstract statement

and smart assistants have improved significantly over the

years (2)mapthatinferencetoaconcretcourseofactionappropriate with respect to speech

recognition [25, 24] and user satis- for the context (e.g., changing the settings of specific

devices). faction [20, 16], a central challenge remains: how can these In this paper, we

leverage the observation that recent

task- assistantsbemade torespondappropriatelytoambiguoususer agnostic large language

models (LLMs) like GPT-3 embody

a commandsthatmaybeinfluencedbycontextorareotherwise vastamountofcross-domain,sometimesunpredictab

that existing rule-based home assistant systems lack,

impossibleforsystemdeveloperstoanticipatebeforehand?An which can make them powerful

tools for inferring user intent example of such a command might be a user preparing their and

generating appropriate context-dependent responses during home to entertain for guests, who

asks their smart assistant to smart home interactions. We first explore the feasibility of a ?get

ready for a party?. The hope is that the assistant?if it is system that places an LLM at the center

of command inference truly smart? might be able to help by inferring the meaning and action

planning, showing that LLMs have the capacity

to infer intent behind vague, context-dependent commands like “get of the statement and

determining how to change the state

of ready for a party? and respond with concrete, machine-parseable available devices in response:

perhaps to start up the user’s instructions that can be used to control smart devices. We

party playlist on a smart speaker and change their smart lights further more demonstrate a proof-of-concept implementation

a festive color scheme. In practice, however, such a

request puts an LLM in control of real devices, showing its ability to infer is beyond the capacity of current smart homes

and change device state appropriately with no

fine-tuning or task-specific training. Our work hints at the promise of LLM- Home will sadly

admit: “I’m sorry, I didn’t understand.” driven systems for context-awareness in smart

environments, In this paper, we are motivated by the observation that large motivating future

research in this

area. language models (LLMs) like OpenAI's GPT-3 [3] have shown an impressive ability to generalize to new tasks

performance, as well as the capacity to infer meaning Index Terms? human-centered

computing, artificial intelligence behind semantically complex or abstract statements [15].

We hence, internet of things thus ask the question: can this powerful capacity for

cross-domain contextual reasoning be applied to practical smart tasks, such as text

classification and sentiment analysis. In home applications?

the same year, OpenAI proposed GPT (Generative Pre-trained Transformer) To explore this question, we carry out a feasibility study

Transformer) [22]. Both models use a transformer architecture places GPT-3 in control of a

smart home. We evaluate GPT-

ture [27] that was pre-trained on a massive corpus of text data, 3? s ability to provide high-quality responses to user commands

including books, articles, and websites. The resulting models of varying ambiguity given

only a simple prompt and a data demonstrate impressive results on a wide range of

natural structure containing information about devices that it can control - language processing tasks,

including language translation, text-to-image generation, and the ability to translate natural

language meaning behind ambiguous user commands like "get ready for descriptions into program

implementations. a party?" or "I am tired and I want to sleep?" and respond with "Following the

success of the transformer-based model,

sub-properly-formatted data describing courses of action, enabling sequent studies have

explored ways to improve and expand more intuitive control of smart devices. We

furthermore build the model's performance. In 2019, Radford et al., published a

proof-of-concept implementation that puts GPT-3 in control an updated version of GPT and

called GPT-2 [23]. Building of real devices, showing LLM-driven command inference and

on the success of GPT-2, Brown et al. released GPT-3 in 2020 action planning can function in practice with no fine-tuning

[4]. After that, in 2023, GPT-4 was introduced. It is currently task-specific training required.

Motivated by our results, we use one of the largest and most powerful language models,

and propose future work that can further leverage the power of

more than 1 trillion parameters [18]. At the time of writing, access to LLMs toward building smarter

smart home applications is limited. We therefore base our study on GPT-3. Our key

contributions are as follows: Two popular approaches exist for adapting task-agnostic ?

An experimental setup and study results that show LLMs to new applications: prompt

engineering and fine-tuning. Prompt

engineering refers to the process of designing a prompt and, in response,

quickly a task-specific prompt or template that guides the model to and appropriately change

the state of the smart devices produce relevant outputs for a particular task [29].

These are available in the home, with no task-specific training. prompts generally contain

instructions to the model written in natural language. An implementation that puts a GPT model in control of in

natural language?e.g., ?explain the following passage about smart devices, showing that it can

intuitively respond to a of text?. Fine-tuning, on the other hand, involves

directly varietyofcommands.Whentoldto?setupforaparty?,it training the model on a new task by

providing task-specific responds by turning on a stereo and configuring a group

examples[23].Thekeyadvantageofpromptengineeringover of Hue lights to loop through a

festive set of colors; fine-tuning is that it does not require task-specific

data? giventhecommand?I?mleaving?,itturnsoffallavailable we therefore adopt that approach

here. Within the realm of devices. We trigger these actions by inputting the LLM?s prompt

engineering, there are zero-shot and few-shot learning response directly into smart device

APIs. approaches. Zero-shot approaches provide the model with a ? Analytical results that

suggest responses are variable in single prompt containing instructions and task-specific

infor- quality, dependent on both the devices available and the mation; few-shot approaches

provide examples to the model nature of the user?s command. In essence, further system of

correct input/output pairs. We focus on zero-shot learning. design is necessary to manage the

LLM's tendency to Context-aware spaces leverage sensor information, user ?not know what it doesn't know? in order to produce data (including past behaviors and preferences), and de- consistently high-quality responses. vice state to influence system actions toward meeting user needs [2]. The notion of ?context-awareness? in this sense The following describes the structure of this paper. Sec- has roots in research on ambient intelligence [5]?that is, the tion II situates our work with related research. Section III development of built environments that sense and adapt to describes the experimental setup that we use to demonstrate users. A concrete example of this concept is a home that lever- the feasibility of LLMs as smart home controllers. Section IV ages contextual information to improve energy efficiency [11, presents the results of our exploratory study, while Section V 7]. In an early paper, Yamazaki suggested that smart homes demonstrates a proof-of-concept implementation. Section VI should go beyond automation and instead integrate expressive offers avenues for future work. Section VII concludes. interfaces between the user



and system [28], a goal that is II. BACKGROUND&RELATEDWORK partially realized in smart assistants [20], but with limited This section provides a high-level overview of LLMs and ability to adapt to more complex user commands [13].

Ample theirapplicationsbeforesituatingourworkwithrelatedefforts prior work has approached the issue of context-awareness us- in context-aware smart spaces.

ingtask-specificmodels[21,12,17,14].Whilethesemethods Large language models (LLMs)

have gained significant can achieve high performance given ample task-specific

data, attention in recent years due to their impressive performance

webelievethatthehighzero-shotperformanceofLLMscould on a wide range of natural language processing tasks. In hint at better generalizability without a need for training

data. 2018,DevlinproposedBERT,alanguagerepresentationmodel However, we are aware of

no work to-date that has explored that uses Bidirectional Encoder Representations from

Trans- theuseoftask-agnosticLLMsfordeepercontextualreasoning formers [6] and can be

fine-tuned for a variety of NLP in smart environments. This motivates our feasibility

study.

### III. SYSTEM DESIGN

In this section, we introduce the system design that we use to

explore the feasibility of LLM-driven smart home control. We first assume the use of an

LLM like GPT-3 that provides responses to user prompts written in natural language.

These LLM models are not task-specific, rather, they are trained on an immense amount of

cross-domain textual information and, depending on the structure of the prompt, can

provide Fig. 1: Data structures for expressing smart home device

and responses suited to a variety of different use cases (e.g., writing user context in prompts to an

LLM, a poem, writing code in response to a high-level program description, etc.). We opt to

adapt the model's outputs to our task using zero-shot learning through prompt engineering.

color values. This overall structure is depicted in Fig. 1 and Our challenge is therefore to

package relevant context and illustrated by the example in the following: user commands into

a concise prompt issued to the model, { such that its responses include concrete,

machine-parseable "user":

```
{ change to device state that can be passed off to the appropriate "location": "living_room" smart
```

device APIs. Qualitatively, we want these courses

of } action to be shaped by the model successfully inferring (1) the } intent behind the user

command and (2) the manner in

which the state of available devices can be changed to meet the user's { intent. To that end, we first

define an abstract schema for "devices": { capturing smart home context before describing a

method for "bedroom": { engineering prompts to conversational LLMs that elicit useful, "lights":

```
{ actionable responses. "bedside_lamp": { "state": "off" A. Context Schema
```

```
} } In order for the model to know what actions are available }, to it, we need to package the available devices, their state,
```

```
{ other relevant information? i.e., the context? into a machine- "lights":
```

```
{ parseable format. This package effectively describes the action "overhead": { space available to
```

the model: the knobs it can turn, and "state": "on" information (e.g., which room the user is

in) that might }, influence how it turns them. It also provides a hint about "lamp": { how the

model should format its response. Representations "state": "off" of context can be complex

and have been explored in the } literature [9, 1]. Since our goal is to conduct an

exploratory }, study rather than design an end-to-end solution, we use a "tvs":

{ schemathatisimplebutadequateforourexperimentalsetup. "living\_room\_tv": { We choose

JSON for structuring this data since it is the de- "state":

"off", factodatainterchangeformatforRESTfulAPIusedbymany "volume": 20 smart home

devices [19, 8, 10]. Leveraging a common format } is also advantageous since there is a high

likelihood that the } LLM has been trained on source material that contains it, } which

benefits the model's ability to converse in

it. } Atthetoplevel,contextisacollectionof?key,value?pairs. } There are two relevant contexts:

?user? context that contains immutableinformationabouttheuser?sstate,e.g.,whichroom

Inthisexample,theuser?shomehastworooms?abedroom theyarein,and?device?context,whichcontainsmutablea

and living room? and the user is currently located in the living room. Immutable information about the

devices in the home. Each room. The bedroom has one light turned off, and the

living room has two lights (one

turned on) and a television. In the home, and within each room we define collections of B.

Prompt Engineering devices organized by type, e.g., ?lights?, ?tvs?, etc. Within a collection of

devices, we define each individual device as Having developed a structure for storing

context, we now have a collection of properties about that device, e.g., for a light move to the

practical challenge of engineering prompts that we can define its ?state? property and its ?r?, ?g?,

and ?b? elicit useful responses from the model. able to relate the meaning of ?party? to the devices

available, as well as alter their specific settings in desirable ways. In the next section, we use

this system design to perform qualitative analysis of the model's responses. IV.

**EVALUATION** This section describes the results of our feasibility

study using the experimental setup described in the previous section. Our evaluations address two

high-level questions: 1) How good are the agent's responses? We measure the quality of the agent's responses, in the sense that they include courses of action that can reasonably be thought to meet the user's request and can be easily machine-parsed and executed. 2) How timely are the agent's responses? We also measure the round-trip response latency. This hints at how feasible a practical system is with respect to user experience and responsiveness. To better understand the system from these two perspectives, we design scenarios of increasing complexity and ambiguity Fig. 2: An example prompt and response from ChatGPT, of context and command. This captures the intuition that demonstrating its ability to change device state in response to (1) different smart homes can have different complexity of ambiguous user commands like "get ready for a party". JSON context, from an apartment with a few smart lights to a large is omitted from this figure in favor of a visual depiction. home with many devices and (2) different user commands can have different levels of ambiguity, from direct commands

like ?turnonthelight?towhollyambiguousstatementslike?Iam Our prompts consist of four

segments, as follows: tired?. Evaluating agent responses under these circumstances ?

Framing. Thisportionofthepromptprovidesdirectionto allows us to identify the failure modes

of LLM-driven smart theconversationalagentaboutitsroleintheinteraction? home control given

increasingly challenging prompts with itisbeingaskedtomakedecisionsasanAIthatcontrols

respect to both the context and the nature of the command. a smart home. We open with the

phrase ?You are an AI We use three contexts of increasing complexity, as follows: that

controls a smart home.? ? Simple: Describes a home with a bedroom and living ? Context. This

informs the agent of the user context and room that have one and two lights, respectively,

all devices available in the environment, which scopes the initially off. Lights can either be

on or off but have no spaceofitsactionsandprovidesahintastothestructure other state (e.g.

color). of our desired response. We continue the prompt: ?Here ? Medium: Same as above, but

adds red, green, and blue is the state of the devices in the home, in JSON format: color state

to each of the lights, with expected values in {devices} Here is information about the user:

{user}?, the range [0, 255]. where both contexts are formatted as shown earlier. ? Complex:

Same as above, but adds a TV with on/off ? Command. This portion inserts the user command

and and volume state to the bedroom, as well as a TV and directs the agent to manipulate the

state of devices in smart speaker to the living room (each also with on/off response, as

follows: ?The user issues the command: and volume

state). {command}.Changethedevicestateasappropriate.?The

Eachofthesecontextsisexpressedintheschemadescribed command is written in natural

language, as a user might in Section III. We combine these contexts with three user utter to

their smart assistant. prompts of increasing ambiguity, as follows: ? Formatting. We close the

prompt by requesting the ? Direct: ?Turn on the light.? This command is simple response in

JSON format so that it can be easily parsed since it directly expresses a state change, as well

as and input to a relevant smart device API: ?Provide your a relevant device. Existing home



assistants can easily respond in JSON format. respond to this type of command. An example

prompt with this structure and the correspond- ?

Indirect: ?Get ready for a party.? This command is more in response from ChatGPT 3.5 are depicted

in Figure 2. We ambiguous since it expresses a desired state change,

but can see that by using the proposed context structure inside the

provides no information about which devices are relevant. the prompt, we are able to elicit

responses from the model ? Ambiguous: ?I am tired.? This command is completely that contain

changes to the underlying JSON that accurately ambiguous since it expresses neither a state

change, nor reflect what a user's intent might be. In essence, GPT-3.5 is which devices might

be relevant. We run our tests with each possible combination of these and a subjective

command like ?get ready for a party?, the three contexts and commands (9 total), each for 10

trials. We LLM simply makes up a response? specifically, it turns on save the agent's response

for each trial in a human-readable all the lights in the house, to include the bedroom.

When format, then perform manual rating to measure the quality of we add a speaker and a

television to the context, the model the responses. Our process for rating the quality of

responses now has more relevant knobs to turn, and produces a higher is based on a binary

label, where each is assigned one of the quality response. The tendency to make something

up when following labels based on its quality: the answer is unknown or requires more

context is an open problem and motivates the development of application-specific ?

Poor(0): ? The changes to the devices do not at all reflect methods to mediate between the user and

the model. the intent behind the user command, or the response is For the most ambiguous

command (?I am tired?), we note malformed/garbled. ? that the model delivers poor responses

regardless of context. ? Good(1): ? The changes to the devices are reasonable for In all but a few

cases, the LLM simply turns on all of the the command. You can imagine someone being

satisfied home ? slight. The exception is in a Medium/Ambiguous trial, with the result, even if it is

somewhat subjective (e.g., where it only turns on the bedroom lamp, perhaps to help

the based on different personal preferences).? user prepare for bed. This is to be expected: an individual?s Three researchers independently reviewed all responses and intent and preference in this case are highly subjective (are assigned them a label. We report the aggregate score for each they, e.g., tired and ready for bed or tired but they have a trial as the average across all assigned scores. We also note the pressing deadline?) and the LLM ultimately cannot read the average latency for each trial?this includes both the network user?s mind. However, since the LLM does not ?know what its transmission time of the request, as well as the inference time doesn?t know?, it does not ask for clarification or go with the taken by the model. Since this time is subject to network safety choice, which is likely to do nothing. Instead, it makes conditions and API demand, it should be taken as a rough something up. estimate rather than a concrete benchmark. Our results are Since our previous results suggested more context is often summarized in Table

I. beneficial, we dig deep to see if we can help it make a better Response time is a function of

context complexity. With choice. We amend the vague command ?I am tired? to offer respect to latency, we can see that responses generally arrive hints at the user?s intent: on the order of seconds, meaning that a practical system ? Ambiguous\*: ?I am tired and I need to work.? could feasibly leverage an LLM for ambiguous command ? Ambiguous\*\*: ?I am tired and I want to sleep.? inference and action planning without significant detriments Provided this added hint at the user?s context, the response to user experience or responsiveness. For direct commands, a quality improves significantly. For Ambiguous\*, the model 2 to 3 second response time may be too long?future system consistently responds by turning on the living room lights designs could thus leverage the LLM only for commands that while leaving all other devices off; for Ambiguous\*\*, the require it. This may entail a hybrid of rule-based inference model turns on only the user?s bedside lamp and, in some for common commands, along with LLM inference for less cases, reduces the volume on their speaker and TV. Note that familiar commands. It is also worth noting that as the context although the amended

statement includes additional context, increases in complexity, the response latency also

increases? it still requires the model to infer meaning in a way that more this motivates future

work to develop methods for filtering rigid or rule-based approaches cannot. context prior to

prompting the agent so that only the most relevant information is provided. V.

**IMPLEMENTATION** Response quality is a function of context and command ambiguity.

With respect to response quality, we find that the To demonstrate LLM-driven smart home

control in prac- LLMapproachprovidesgoodresponsesgiventhesamedirect tice, we built a

proof-of-concept implementation in Python. and simple commands that current home

assistants are able to Our implementation accepts user commands as strings, pack- service.

Note, however, that unlike existing home assistants, ages them into prompts along with

contextual information theLLMapproachutilizesamuchsimplesystemarchitecture about real

devices, then processes responses from OpenAI?s that performs command inference and

action planning in text-davinci-003 model into smart home API calls that the same pass.

These results are consistent given increasing change the device state as specified by the

LLM. We scope degrees of context complexity, suggesting that the model was the

application to one room (a researcher's living room) with not overwhelmed by the growth in

the decision space that three Philips Hue color smart lights [19] and one

TP-Link comes with adding new devices and possible state changes. On Kasa smart plug [26] that

controls a stereo. We store the the contrary, the model provides better responses when given

device context in JSON as in our experimental setup, with more context that might be

relevant to the user's command, the difference that the device state for the Hue lights is

pulled as is apparent when comparing the low response quality of directly from the Hue API

without modification. Our code is the Simple/Indirect and Medium/Indirect experiments

against open source and available online.<sup>1</sup> the Complex/Indirect experiment. Upon inspection

of the responses, the reason for this is clear: given minimal context

<sup>1</sup><https://github.com/UT-MPC/homegpt> Context Command Avg. Quality

AvgLatency(sec) context will be necessary. This will also become essential Simple Direct

1.00 2.42 Indirect 0.67 2.31 as the amount of context grows to include sensor data,

user Ambiguous 0.00 2.22 preference data, and a growing and more diverse set of Medium

Direct 1.00 4.56 controllable devices. We note that in our experiments, we did Indirect 0.63

4.70 not attempt to test the limits of how much context a model Ambiguous 0.17

4.97 can receive before the quality or latency of responses degrades Complex Direct 1.00

7.90 Indirect 1.00 7.25 substantially. This should be considered in future work. Ambiguous

0.00 7.04 Robust system design. While we were able to leverage Complex Ambiguous\* 1.00

7.49 a simple system design in this paper, an end-to-end system Ambiguous\*\* 1.00 8.09 will

need a more robust design to account for several factors. TABLE I: Results for experiments

given various combina- First, since LLMs do not yet ?know what they don?t

know?, tionsofdifferentcontextcomplexityandcommandambiguity. the likelihood of invalid or

low-quality responses remains Higherqualityresponses suggest the model produced a course

high. In the case of responses where the model makes

invalid of action that would be desirable for an end user (e.g., turning changes to device state (e.g, to

add new settings to a device), on the bedroom light when receiving the command "I am tired a full

system should include a way to enforce a set of formal and I want to sleep?"). Lower latency

suggests better system properties for device states. In the case of

unsatisfactory responsiveness. responses, it would be beneficial to develop a method

for learning user preferences or seeking clarifying information (e.g., "are you tired and want to

sleep, or are you tired but The teaser figure depicts the result when issuing the com- need an

energy boost?"). mand "set up for a party?". We include the JSON context of From commands to

automation. Our primary focus in the light group2, along with a field for the plug powering

this exploratory study was on immediate commands? the user the stereo. The model mutates the

parameters in the JSON to makes a request and the model immediately responds

with change the stereo state to "on" and, impressively, also changes a state change. Future work



could investigate the use of the "effect" parameter of the Hue light group from "none" to LLMs

for more intuitive automation planning. A user could, "colorloop" to create a looping color

effect. The latter change for instance, ask their smart assistant to "play jazz when it suggests

that GPT-3 may have been trained on material about rains" and the model could leverage

contextual information to the specific features of the Hue API and can leverage that put in

place an automation sequence that meets their needs. along with the inferred meaning of the

user command to This would obviate the need for pre-programmed automation trigger more

intuitive changes than existing systems.

routines and could substantially improve user satisfaction with We briefly list multiple other

commands we tested in our smart assistant systems. implementation, along with responses

from the model: VII. CONCLUSION "make it bright in here" sets lights to full brightness In

this paper, we explored the feasibility of smarter smart "make it groovy" sets lights to color

loop; adds invalid home control using large language models (LLMs). We pro- "genre" field to

stereo and sets it to "groovy" posed a simple system design for capturing smart home con-

"gotta relax" dims lights, turns stereo on text (i.e., information about the user and controllable

devices "I'm cold" sets lights to warm white, turns stereo

on in the environment) in engineered prompts to GPT-3, showing "I'm leaving" turns off lights and

stereo that the model has the ability to infer meaning behind indirect "I'm home" turns on lights and

stereo and ambiguous user commands like "I am tired and I need to" We note, of course, that

these tests are far from

exhaustive. work and, in response, generate changes to smart device state. We observed high

variability in responses, meaning the same We implemented our system design, giving

GPT-3 control command can elicit many responses: some good, some bad. of real devices

and finding that it is able to quickly and A more robust system design will be necessary to

tackle the appropriately control them in response to user commands with inconsistencies present

in current LLM model outputs. We no fine tuning and no post-processing of its responses.

By address this in our discussion in the following section. simply telling GPT-3 what devices are available and what the VI. DISCUSSION&FUTUREWORK user wants, it can generate courses of action in response. Our work hints at the capability of GPT-3 and similar Our efforts in this paper hint at exciting opportunities for models to go far beyond the current abilities of smart space future work. We suggest several avenues for further research. controlandmotivatesfutureworkwithcontextmodeling,end- Managing contextual information. We found that includ- to-end system design, and approaches for further leveraging ing more context can improve the quality of the model?s GPT-3?s capabilites to develop complex automation routines responses, but at the expense of response latency. To effec- in response to user commands. tively navigate this tradeoff in an end-to-end solution, a more involved approach for storing, pre-processing, and expressing 2<https://developers.meethue.com/develop/hue-api/groups-api/>REFERENCES

[25] Prakhar Swarup et al. ?Improving ASR confidence scores for

Alexa using acoustic and hypothesis embeddings?. In: (2019). [1] Hamim Md Adal et al. ?The

Space Broker: A Middleware for Mediating Interactions in Smart IoT Spaces?. In:

Proceedings of the [26] TP-Link Kasa Smart Plugs. 2023. URL:

<https://www.kasasmart.com/> 8th ACM International Conference on Systems for

Energy-Efficient us/products/smart-plugs. Buildings, Cities, and Transportation. BuildSys ?21.

Coimbra, Portu- [27] Ashish Vaswani et al. ?Attention is all you need?. In: Advances

in neural information processing systems 30 (2017). gal: Association for Computing Machinery, 2021, pp. 101?110.

Tatsuya Yamazaki. ?Beyond the Smart Home?. In: 2006 International 9781450391146. DOI:

10.1145/3486611.3486664. URL:

<https://doi.org/10.1145/3486611.3486664> Conference on Hybrid Information Technology. Vol. 2. 2006, pp. 350?359.

Muhammad Raisul Alam, Mamun Bin Ibne Reaz, and Mohd Alauddin [29] Tony Z. Zhao et al.

Calibrate Before Use: Improving Few-Shot Per- Mohd Ali. ?A Review of Smart Homes? Past,

Present, and Future?. Performance of Language Models. 2021. arXiv:2102.09690[cs.CL]. In: IEEE

Transactions on Systems, Man, and Cybernetics, Part

C (Applications and Reviews) 42.6(2012), pp. 1190–1203. DOI: 10.1109/ TSMCC.2012.2189204. [3]

Tom Brown et al. "Language models are few-shot learners". In:

Advances in neural information processing systems 33(2020), pp. 1877–1901. [4]

Tom B. Brown et al. "Language Models are Few-Shot Learners". In: CoRR abs/2005.14165(2020). arXiv:2005.14165. [5]

URL: <https://arxiv.org/abs/2005.14165>. [5] Diane J Cook, Juan C Augusto, and Vikramaditya

R Jakkula. "Ambient intelligence: Technologies, applications, and opportunities".

In: Pervasive and mobile computing 5.4(2009), pp. 277–298. [6]

Jacob Devlin et al. "Bert: Pre-training of deep bidirectional transformers for language understanding". In: arXiv preprint

PRGerald Filho et al. "Energy-efficient smart homes systems: Infrastructure and decision-making process". In: International

Google Nest API Documentation. 2022. URL:

<https://developers.google.com/nest/device-access>. [9] Jie Hua et al. "CoPI: Enabling

Probabilistic Conflict Prediction in Smart Space Through Context-awareness". In: 2022

IEEE/ACM Seventh International Conference on Internet-of-Things Design

and Implementation(IoTDI).IEEE.2022,pp.30-42. [10] Insteon REST API Documentation.

2022. URL: <https://www.insteon.com/developer>. [11] Marco Jahn et al. "The energy aware

smart home?". In: 2010

5th international conference on future information technology. IEEE. 2010, pp.1-8. [12] M

Humayun Kabir et al. "Machine learning based adaptive

context- aware system for smart home environment?". In: International Journal of Smart Home 9.11(2015), pp.55-62

Julia Kiseleva et al. "Understanding user satisfaction with

intelligent assistants?". In: Proceedings of the 2016 ACM on Conference on Human Information Interaction and Retrieval

Tiankai Lian et al. "An unsupervised user behavior prediction algorithm based on machine learning and neural network

Xiao Liu et al. "GPT understands, too?". In: arXiv preprint arXiv:2103.10385(2021). [16]

Irene Lopatovska et al. "Talk to me: Exploring user interactions with the Amazon Alexa?". In: Journal

of Librarianship and Information Science 51.4(2019), pp.984-997. [17] Isaac Machorro-Cano et

al. ?HEMS-IoT: A big data and machine learning-based smart home system for energy saving?.

In: *Energies* 13.5(2020),p.1097. [18]

OpenAI.GPT-4.URL:<https://openai.com/product/gpt-4>. [19] Philips Hue API Documentation.

2022. URL: <https://developers.meethue.com/>. [20]

AmandaPuringtonetal.??AlexaismynewBFF?SocialRoles,User Satisfaction,andPersonificationoftheAmazon

the 2017 CHI conference extended abstracts on human factors

in computingsystems.2017,pp.2853?2859. [21] Basheer Qolomany et al. ?Leveraging machine

learning and big data for smart buildings: A comprehensive survey?. In: *IEEE Access*

7 (2019),pp.90316?90356. [22]

AlecRadfordetal.?Improvinglanguageunderstandingbygenerative pre-training?.In:(2018). [23]

Alec Radford et al. ?Language models are unsupervised

multitask learners?.In:OpenAIblog1.8(2019),p.9. [24] Anirudh Raju et al. ?Scalable multi

corpora neural language models forasr?.In:arXivpreprintarXiv:1907.01677 (2019).